

Compsci 340 Assignment 3

PART 1

1.1

The program header contains the segments (in the address space of a process running that ELF executable) projected in virtual memory at exec time. The final sections header defines the linkable point of view. Each section belongs to a segment and may or may not be visible – i.e. mapped into memory – at execution time). The ELF file header tells where program header table & section header table are.

1.2

```
rahul@rahul-VirtualBox:~/Desktop/A3$ gcc hello.c -o hello
rahul@rahul-VirtualBox:~/Desktop/A3$ ./hello
Hello riss899!
00400000-00401000 r-xp 00000000 08:01 550955 /home/rahul/Desktop/A3/hello
00600000-00601000 r--p 00000000 08:01 550955 /home/rahul/Desktop/A3/hello
00601000-00602000 rw-p 00001000 08:01 550955 /home/rahul/Desktop/A3/hello
006d3000-006f4000 rw-p 00000000 00:00 0 [heap]
7fcff3436000-7fcff35f6000 r-xp 00000000 08:01 6168 /lib/x86_64-linux-gnu/libc-2.23.so
7fcff35f6000-7fcff37f6000 ---p 001c0000 08:01 6168 /lib/x86_64-linux-gnu/libc-2.23.so
7fcff37f6000-7fcff37fa000 r--p 001c0000 08:01 6168 /lib/x86_64-linux-gnu/libc-2.23.so
7fcff37fa000-7fcff37fc000 rw-p 001c4000 08:01 6168 /lib/x86_64-linux-gnu/libc-2.23.so
7fcff37fc000-7fcff3800000 rw-p 00000000 00:00 0
7fcff3800000-7fcff3826000 r-xp 00000000 08:01 6140 /lib/x86_64-linux-gnu/ld-2.23.so
7fcff3a0a000-7fcff3a0d000 rw-p 00000000 00:00 0
7fcff3a23000-7fcff3a25000 rw-p 00000000 00:00 0
7fcff3a25000-7fcff3a26000 r--p 00025000 08:01 6140 /lib/x86_64-linux-gnu/ld-2.23.so
7fcff3a26000-7fcff3a27000 rw-p 00026000 08:01 6140 /lib/x86_64-linux-gnu/ld-2.23.so
7fcff3a27000-7fcff3a28000 rw-p 00000000 00:00 0
7ffdf186c000-7ffdf188d000 rw-p 00000000 00:00 0 [stack]
7ffdf18ff000-7ffdf1901000 r--p 00000000 00:00 0 [vvar]
7ffdf1901000-7ffdf1903000 r-xp 00000000 00:00 0 [vdso]
fffffffff600000-fffffffff601000 r-xp 00000000 00:00 0 [vsyscall]
0 1 2 3 4 5 6 7 8 9 Goodbye world.
rahul@rahul-VirtualBox:~/Desktop/A3$
```

1.3

- 1) When the operating system loads the resulting file, only the allocable part is mapped into memory. The non-allocable part remains in the file, but is not visible in memory.
- 2) Since you have addresses bigger than 32 bits, some of the shared objects have “/lib/x86_64-linux-gnu/” in their path. Therefore, its amd64.

1.4

```
rahul@rahul-VirtualBox:~/Desktop/A3$  
rahul@rahul-VirtualBox:~/Desktop/A3$ nm hello  
000000000060106f B __bss_start  
0000000000601060 D bye  
0000000000601070 b completed.7585  
0000000000601050 D __data_start  
0000000000601050 W data_start  
00000000004005e0 t deregister_tm_clones  
0000000000400660 t __do_global_dtors_aux  
0000000000600e18 t __do_global_dtors_aux_fini_array_entry  
0000000000601058 D __dso_handle  
0000000000600e28 d _DYNAMIC  
000000000060106f D _edata  
0000000000601078 B _end  
00000000004007d4 T _fini  
0000000000400680 t frame_dummy  
0000000000600e10 t __frame_dummy_init_array_entry  
0000000000400958 r __FRAME_END__  
U getpid@@GLIBC_2.2.5  
0000000000601000 d _GLOBAL_OFFSET_TABLE_  
w __gmon_start__  
000000000040080c r __GNU_EH_FRAME_HDR  
0000000000400500 T _init  
0000000000600e18 t __init_array_end  
0000000000600e10 t __init_array_start  
00000000004007e0 R _IO_stdin_used  
w _ITM_deregisterTMCloneTable  
w _ITM_registerTMCloneTable  
0000000000600e20 d __JCR_END__  
0000000000600e20 d __JCR_LIST__  
w _Jv_RegisterClasses  
00000000004007d0 T __libc_csu_fini  
0000000000400760 T __libc_csu_init  
U __libc_start_main@@GLIBC_2.2.5  
00000000004006d8 T main  
00000000004006a6 T numbers  
U printf@@GLIBC_2.2.5  
U puts@@GLIBC_2.2.5  
0000000000400620 t register_tm_clones  
U sprintf@@GLIBC_2.2.5  
U __stack_chk_fail@@GLIBC_2.4  
00000000004005b0 T _start  
U system@@GLIBC_2.2.5  
0000000000601070 D __TMC_END__  
0000000000601074 B unused  
rahul@rahul-VirtualBox:~/Desktop/A3$
```

- Stdio.h
- sys/types

- `unistd.h`
- `stdlib.h`
- `bye[]`

The symbol is in the initialised data section

- `int unused`
The symbol is in the uninitialized data section (known as BSS).
- `void numbers()`

The symbol is in the text (code) section.

- `puts("Hello riss899!");`

The symbol is undefined

- `numbers();`
The symbol is in the text (code) section.

- `sprintf(mem, "cat /proc/%d/maps", pid);`
The symbol is undefined

- `system(mem);`

The symbol is undefined

- `puts(bye);`

The symbol is undefined

Local variables are not accessible from object files, `nm` lists the symbol table of object files, symbol tables are used for linking, local variables are not needed at link time.

1.5

- A text file is a kind of computer file that is structured as a sequence of lines of electronic text. A text file exists stored as data within a computer file system.

- In computing, a data segment (often denoted .data) is a portion of an object file or the corresponding virtual address space of a program that contains initialized static variables, that is, global variables and static local variables. The data segment is read-write, since the values of variables can be altered at run time. This contrasts with the read-only data segment (rodata segment or .rodata), which contains static constants rather than variables.
- The name .bss or bss is used by many compilers and linkers for a part of the data segment containing statically-allocated variables represented solely by zero-valued bits initially. It is often referred to as the “bss section” or “bss segment”. The BSS segment contains all global variables and static variables that are initialized to zero or do not have explicit initialization in source code.

```

rahul@rahul-VirtualBox:~/Desktop/A3$ readelf -lW hello
Elf file type is EXEC (Executable file)
Entry point 0x4005b0
There are 9 program headers, starting at offset 64

Program Headers:
Type           Offset             VirtAddr           PhysAddr          FileSiz  MemSiz   Flg  Align
PHDR           0x000040          0x0000000000400040 0x0000000000400040 0x0001f8 0x0001f8 R E  0x8
INTERP         0x000238          0x0000000000400238 0x0000000000400238 0x00001c 0x00001c R   0x1
              [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]
LOAD           0x000000          0x0000000000400000 0x0000000000400000 0x00095c 0x00095c R E  0x200000
LOAD           0x000e10          0x0000000000600e10 0x0000000000600e10 0x00025f 0x000268 RW  0x200000
DYNAMIC         0x000e28          0x0000000000600e28 0x0000000000600e28 0x0001d0 0x0001d0 RW  0x8
NOTE           0x000254          0x0000000000400254 0x0000000000400254 0x000044 0x000044 R   0x4
GNU_EH_FRAME   0x00080c          0x000000000040080c 0x000000000040080c 0x00003c 0x00003c R   0x4
GNU_STACK      0x000000          0x0000000000000000 0x0000000000000000 0x000000 0x000000 RW  0x10
GNU_RELRO      0x000e10          0x0000000000600e10 0x0000000000600e10 0x0001f0 0x0001f0 R   0x1

Section to Segment mapping:
Segment Sections...
00
01      .interp
02      .interp.note.ABI-tag.note.gnu.build-id.gnu.hash.dynsym.dynstr.gnu.version.gnu.version_r.rela.dyn.rela.plt.init.plt.plt.got.text
.fini.rodata.eh_frame_hdr.eh_frame
03      .init_array.fini_array.jcr.dynamic.got.got.plt.data.bss
04      .dynamic
05      .note.ABI-tag.note.gnu.build-id
06      .eh_frame_hdr
07
08      .init_array.fini_array.jcr.dynamic.got

```

1.6

- The last digits of the offset are the same of the last digits of the Virtual address

1.7

Seg	Type	Virtual address	Size	Access(flag)	Map area (range of addresses)	File offset	Mem access
2	LOAD	0x000000400000	0x00095c	R E	00400000- 00401000	0x000000	r-xp
3	LOAD	0x000000600e10	0x000268	RW	00601000- 0060200	0x000e10	r-wp

1.8

.text – is executable so read /
exec

.rodata – is read only so read,

.data – read and write so rw

.bss – read and write so rw

```
2 .note.gnu.build-id 00000024 0000000000400274 0000000000400274 00000274 2**2
CONTENTS, ALLOC, LOAD, READONLY, DATA
3 .gnu.hash 00000040 0000000000400298 0000000000400298 00000298 2**3
CONTENTS, ALLOC, LOAD, READONLY, DATA
4 .dynsym 00000750 00000000004002d0 00000000004002d0 000002d0 2**3
CONTENTS, ALLOC, LOAD, READONLY, DATA
5 .dynstr 00000316 0000000000400a28 0000000000400a28 00000a28 2**0
CONTENTS, ALLOC, LOAD, READONLY, DATA
6 .gnu.version 0000009c 0000000000400d3e 0000000000400d3e 00000d3e 2**1
CONTENTS, ALLOC, LOAD, READONLY, DATA
7 .gnu.version_r 00000060 0000000000400de0 0000000000400de0 00000de0 2**3
CONTENTS, ALLOC, LOAD, READONLY, DATA
8 .rela.dyn 00000090 0000000000400e40 0000000000400e40 00000e40 2**3
CONTENTS, ALLOC, LOAD, READONLY, DATA
9 .rela.plt 00000078 0000000000400ed0 0000000000400ed0 00000ed0 2**3
CONTENTS, ALLOC, LOAD, READONLY, DATA
10 .init 0000001a 0000000000401548 0000000000401548 00001548 2**2
CONTENTS, ALLOC, LOAD, READONLY, CODE
11 .plt 00000460 0000000000401570 0000000000401570 00001570 2**4
CONTENTS, ALLOC, LOAD, READONLY, CODE
12 .plt.got 00000008 00000000004019d0 00000000004019d0 000019d0 2**3
CONTENTS, ALLOC, LOAD, READONLY, CODE
13 .text 00007369 00000000004019e0 00000000004019e0 000019e0 2**4
CONTENTS, ALLOC, LOAD, READONLY, CODE
14 .fini 00000009 0000000000408d4c 0000000000408d4c 00008d4c 2**2
CONTENTS, ALLOC, LOAD, READONLY, CODE
15 .rodata 00001233 0000000000408d60 0000000000408d60 00008d60 2**5
CONTENTS, ALLOC, LOAD, READONLY, DATA
16 .eh_frame_hdr 00000304 0000000000409f94 0000000000409f94 00009f94 2**2
CONTENTS, ALLOC, LOAD, READONLY, DATA
17 .eh_frame 000010d4 000000000040a298 000000000040a298 0000a298 2**3
CONTENTS, ALLOC, LOAD, READONLY, DATA
18 .init_array 00000008 000000000060be10 000000000060be10 0000be10 2**3
CONTENTS, ALLOC, LOAD, DATA
19 .fini_array 00000008 000000000060be18 000000000060be18 0000be18 2**3
CONTENTS, ALLOC, LOAD, DATA
20 .jcr 00000008 000000000060be20 000000000060be20 0000be20 2**3
CONTENTS, ALLOC, LOAD, DATA
21 .dynamic 000001d0 000000000060be28 000000000060be28 0000be28 2**3
CONTENTS, ALLOC, LOAD, DATA
22 .got 00000008 000000000060bff8 000000000060bff8 0000bff8 2**3
CONTENTS, ALLOC, LOAD, DATA
23 .got.plt 00000240 000000000060c000 000000000060c000 0000c000 2**3
CONTENTS, ALLOC, LOAD, DATA
24 .data 000000b4 000000000060c240 000000000060c240 0000c240 2**5
CONTENTS, ALLOC, LOAD, DATA
25 .bss 000009a8 000000000060c300 000000000060c300 0000c2f4 2**5
ALLOC
26 .gnu_debuglink 00000034 0000000000000000 0000000000000000 0000c2f4 2**0
CONTENTS, READONLY
```

PART 2

2.1

An example of backward compatibility is the Windows 64-bit. It has software called WOW64 that provides compatibility by emulating a 32-bit system. Consequences of moving from 32bit to 64 bit are:

- You need more memory for many operations
- The effective part of processor cash is smaller
- The size of code also increases because of additional prefixes and instructions containing 8-byte operands instead of 4-byte ones.

2.2

- Since each node holds a frame number, it must also point to the next node in the free list. Therefore: $32\text{bits} * 2 = 64\text{bits} = 8\text{ bytes}$
- In Kernel, it will have to store all the 2^{21} bits which is calculated from $2^{34} / 2^{13} = 2^{21}$. When the system is starting, it must know which

memory frames are free, so it stores all of them. This means that 2^{18} bytes are stored.

- Extent is a chunk of storage space logical volume management which is used internally to provide different device mappings. Extent will have to ability to remove the metadata of the larger files.

2.3

page – 4KB

2^{32} of swap space – physical space = max useful to allocate
swap is slow ram is cheap maybe reallocate zero swap space.

2.4

$H = \text{TLB hit ratio} = 99\% = 0.99$

TLB access time = 1ns

Memory access time = 50ns

PT access = 100ns

$(H)(\text{TLB access time} + \text{mem access time}) + (1-H)(\text{TLB access} + \text{PT access} + \text{mem access})$
 $(0.99) * (1 + 50) + (1 - 0.99) * (1 + 100 + 50) = 52\text{ns}$

2.5

- Logical address bit = 64 bits
- $64 - \log_2(2^{22}) = 42$ bits
- Number of page will be $= 2^{64} / 2^{22} = 2^{42}$.
- Pages we have entry of page table = 2^6
- Number of entry in 1 page will be $2^{22} / 2^6 = 2^{16}$
- Bit to represent one entry = $\log_2(2^{16}) = 16$
- And bit for page is 42 bits
- Therefore, number of level = $42 / 16 = 3$ level page table

Since there is three levels plus one page for data/program and one for the stack which makes it 5 for the minimum page number.

2.6

- Normal instruction takes 1 microsecond (10^{-6} sec)
- Instruction with page fault takes 2000 micro seconds

- Given the program takes 60 sec and there were 20000 p.f
- Time taken by page faults = $20000 * 2000 \text{ microseconds} = 40 \text{ sec}$
- Rest of the 20 sec is consumed by program execution
- $2 * 10,000 = 20 \text{ seconds} = \text{the time it takes to deal with page faults with twice as much memory as execution time.}$
- Therefore $20 + 20 = 40 \text{ seconds.}$

2.7

a. FIFO

	1	2	3	4	3	5	6	7	6	5	4	3	4	7	6	1	5	4	1	2
0	1						6													
0		2						7												
0			3													1				
0				4																2
0						5														
PF	F	F	F	F		F	F	F								F				F

PF = 9

b. LRU

	1	2	3	4	3	5	6	7	6	5	4	3	4	7	6	1	5	4	1	2
0	1						6													
0		2						7												2
0			3														5			
0				4																
0						5										1				
PF	F	F	F	F		F	F	F								F	F			F

PF = 10

c. LFU

	1	2	3	4	3	5	6	7	6	5	4	3	4	7	6	1	5	4	1	2
0	1						6													
0		2						7												2
0			3													1				
0				4																
0						5														
PF	F	F	F	F		F	F	F								F				F

PF = 9

d. Optimal

	1	2	3	4	3	5	6	7	6	5	4	3	4	7	6	1	5	4	1	2
0	1							7												
0		2					6													
0			3													1				2
0				4																
0						5														
PF	F	F	F	F		F	F	F								F				F

PF = 9